

Pointers and Variable-Length Arrays (C99)

- Pointers are allowed to point to elements of variable-length arrays (VLAs).
- Example:

```
void f(int n) {
    int a[n], *p;
    p = a;
    ...
}
```
- When the VLA has more than one dimension, the type of the pointer depends on the length of each dimension except for the first.

A two-dimensional example:

```
void f(int m, int n)
{
    int a[m][n], (*p)[n];
    p = a;
    ...
}
```

- Since the type of `p` depends on `n`, which isn't constant, `p` is said to have a ***variably modified type***.
- A potential problem:**
The validity of an assignment such as `p = a` can't always be determined by the compiler.

The following code will compile but is correct only if `m` and `n` are equal:

```
int a[m][n], (*p)[m];
p = a;
```

If `m` is not equal to `n`, any subsequent use of `p` will cause undefined behavior.

Pointer arithmetic works with VLAs.

- A two-dimensional VLA:

```
int a[m][n];
```

A pointer capable of pointing to a row of `a`:

```
int (*p)[n];
```

A loop that clears column `i` of `a`:

```
for(p = a; p < a + m; p++)
    (*p)[i] = 0;
```
- Restriction on variable-length arrays:
the declaration must be inside the body of a function or in a function prototype.

Chapter 13 Strings

- Strings are arrays of characters in which a special character—the null character—marks the end.
- A ***string literal*** (or string constant) is a sequence of characters enclosed within double quotes:
"When you come to a fork in the road, take it."
String literals may contain escape sequences.
- Character escapes often appear in `printf` and `scanf` format strings.
For example, each `\n` character in the string
"Candy\nIs dandy\nBut liquor\nIs quicker.\n--Ogden Nash\n" causes the cursor to advance to the next line:
Candy
Is dandy
But liquor
Is quicker.
--Ogden Nash
- The backslash character (`\`) can be used to continue a string literal from one line to the next:

```
printf("When you come to a fork in the \
road, take it. --Yogi Berra");
```

Another way to do it:

```
printf("When you come to a fork in the "
      "road, take it. --Yogi Berra");
```

How String Literals Are Stored

- When a C compiler encounters a string literal of length `n` in a program, it sets aside `n + 1` bytes of memory for the string.

The memory will contain the characters in the string, plus one extra character—the ***null character***—to mark the end of the string.

The null character is a byte whose bits are all zero, so it's represented by the `\0` escape sequence.

For example, the string literal "abc" is stored as an array of four characters:

a	b	c	\0
---	---	---	----

- The string "" is stored as a single null character:



- Since a string literal is stored as an array, the compiler treats it as a pointer of type `char *`.
- We can use a string literal wherever C allows a `char *` pointer:

```
char *p;
p = "abc";
```

This assignment makes `p` point to the first character of the string.

- String literals can be subscripted:

```
char ch;
ch = "abc"[1];
```

The new value of `ch` will be the letter `b`.

- A function that converts a number between 0 and 15 into the equivalent hex digit:

```
char digit_to_hex_char(int digit){
    return "0123456789ABCDEF"[digit];
}
```

- A string literal containing a single character isn't the same as a character constant.
 - `"a"` is represented by a *pointer*.
 - `'a'` is represented by an *integer*.

- A legal call of `printf`:

```
printf("\n");
```

An illegal call:

```
printf('\n');    /** WRONG **/
```

String Variables

- Any one-dimensional array of characters can be used to store a string.
A string must be terminated by a null character.
- Difficulties with this approach:
 - It can be hard to tell whether an array of characters is being used as a string.
 - String-handling functions must be careful to deal properly with the null character.
 - Finding the length of a string requires searching for the null character.

- If a string variable needs to hold 80 characters, it must be declared to have length 81:

```
#define STR_LEN 80
. . .
char str[STR_LEN+1];
```

Adding 1 to the desired length allows room for the null character at the end of the string.

- Be sure to leave room for the null character when declaring a string variable.

Failing to do so may cause unpredictable results when the program is executed.

- The actual length of a string depends on the position of the terminating null character.

An array of `STR_LEN + 1` characters can hold strings with lengths between 0 and `STR_LEN`.

Initializing a String Variable

- A string variable can be initialized at the same time it's declared:


```
char date1[8] = "June 14";
```

The compiler will automatically add a null character so that `date1` can be used as a string:

- If the initializer is too short to fill the string variable, the compiler adds extra null characters:

```
char date2[9] = "June 14";
```

Appearance of `date2`:



- An initializer for a string variable can't be longer than the variable, but it can be the same length:

```
char date3[7] = "June 14";
```

The code line above is not good.

- Declaration of a string variable may **omit its length**, in which case the compiler computes it:

```
char date4[] = "June 14";
```

The compiler sets aside eight characters for `date4`, enough to store the characters in `"June 14"` plus a null character.

- **Omitting the length of a string variable** is especially useful if the initializer is long, since computing the length by hand is error-prone.

Character Arrays versus Character Pointers

- The declaration

```
char date1[] = "June 14";
```

declares date to be an *array*. The similar-looking

```
char *date2 = "June 14";
```

declares date to be a *pointer*.

However, there are significant differences between the two versions of date.

- In the array version, the characters stored in date1 can be modified. In the pointer version, date2 points to a string literal that shouldn't be modified.
- In the array version, date1 is an array name. In the pointer version, date2 is a variable that can point to other strings.
- The declaration

```
char *p;
```

does not allocate space for a string.
Before we can use p as a string, it must point to an array of characters.
- One possibility is to make p point to a string variable:

```
char str[STR_LEN+1], *p;  
p = str;
```
- Another possibility is to make p point to a dynamically allocated string.
- Using an uninitialized pointer variable as a string is a serious error.

Example: An attempt at building the string "abc":

```
char *p;  
p[0] = 'a';    /** WRONG **/  
p[1] = 'b';    /** WRONG **/  
p[2] = 'c';    /** WRONG **/  
p[3] = '\0';   /** WRONG **/
```

Since p hasn't been initialized, this causes undefined behavior.

Reading and Writing Strings

- **Writing** a string is easy using either printf or puts.

The %s conversion specification allows printf to write a string:

```
char str[]="Are we having fun yet?";  
printf("%s\n", str);
```

The output will be

Are we having fun yet?

- printf writes the characters in a string one by one until it encounters a null character.
- To print part of a string, use the conversion specification %.ps. p is the number of characters to be displayed.
- The statement

```
printf("%.6s\n", str);
```

will print
Are we
- The %ms conversion will display a string in a field of size m.
If the string has fewer than m characters, it will be right-justified within the field.
To force left justification instead, we can put a minus sign in front of m.
- The m and p values can be used in combination.
A conversion specification of the form %m.ps causes the first p characters of a string to be displayed in a field of size m.
- printf isn't the only function that can write strings. The C library also provides **puts**:

```
puts(str);
```

After writing a string, puts always writes an additional new-line character.
- **Reading** a string is a bit harder. To read a string in a single step, we can use either **scanf** or **gets**.
An alternative is to read strings one character at a time.
- The %s conversion specification allows scanf to read a string into a character array:

```
scanf("%s", str);
```

str is treated as a pointer, so there's no need to put the & operator in front of str.
- When scanf is called, it skips white space, then reads characters and stores them in str until it encounters a white-space character.
- scanf always stores a null character at the end of the string.
- scanf won't usually read a full line of input.
A new-line character will cause scanf to stop reading, but so will a space or tab character.

- To read an entire line of input, we can use **gets**.
Properties of gets:
 - Doesn't skip white space before starting to read input.
 - Reads until it finds a new-line character.
 - Discards the new-line character instead of storing it; the null character takes its place.

- Consider the following program fragment:

```
char sentence[SENT_LEN+1];
printf("Enter a sentence:\n");
scanf("%s", sentence);
```

- Suppose that after the prompt

Enter a sentence:

the user enters the line

To C, or not to C.

scanf will store the string "To" in sentence.

- Suppose that we replace **scanf** by **gets**:

```
gets(sentence);
```

When the user enters the same input as before, **gets** will store the string "To C, or not to C." in sentence.

- As they read characters into an array, scanf and gets have **no way to detect when it's full**.

Consequently, they may store characters past the end of the array, causing undefined behavior.

- scanf can be made safer by using the conversion specification **%ns** instead of **%s**.

n is an integer indicating the maximum number of characters to be stored.

- **gets** is inherently unsafe; **fgets** is much better.

char *fgets(char *str, int n, FILE *stream)

reads a line from the specified stream and stores it into the string pointed to by **str**. It stops when either (**n-1**) characters are read, the newline character is read, or the end-of-file is reached, whichever comes first.

- Programmers often write their own input functions.

Issues to consider:

- Should the function skip white space before beginning to store the string?
- What character causes the function to stop reading: a new-line character, any white-space character, or some other character? Is this character stored in the string or discarded?

- What should the function do if the input string is too long to store: discard the extra characters or leave them for the next input operation?

Reading Strings Character by Character

- Suppose we need a function that (1) doesn't skip white-space characters, (2) stops reading at the first new-line character (which isn't stored in the string), and (3) read *n* characters and discards extra characters.

- A prototype for the function:

```
int read_line(char str[], int n);
```

If the input line contains more than *n* characters, **read_line** will discard the additional characters.

read_line will return the number of characters read.

- **read_line** consists primarily of a loop that calls **getchar** to read a character and then stores the character in **str**, provided that there's room left:

```
int read_line(char str[], int n) {
    int ch, i = 0;
    while((ch = getchar()) != '\n')
        if (i < n) str[i++] = ch;
    str[i] = '\0'; //terminates string
    return i; //No. of characters stored
}
```

- **ch** has **int** type rather than **char** type because **getchar** returns an **int** value.
- Before returning, **read_line** puts a null character at the end of the string.

Functions **scanf** and **gets** automatically put a null character at the end of an input string.

If we're writing our own input function, we must take on that responsibility.

Accessing the Characters in a String

- Since strings are arrays, we can use subscripting to access the characters in a string.

To process every character in a string **s**, we can set up a loop that increments a counter **i** and selects characters via the expression **s[i]**.

An example we saw earlier:

```
char digit_to_hex_char(int digit){
    return "0123456789ABCDEF"[digit];
}
```

- Another example: A function that counts the number of spaces in a string:

```
int count_spaces(const char s[]) {
    int count = 0, i;
    for (i = 0; s[i] != '\0'; i++)
        if (s[i] == ' ') count++;
    return count;
}
```

- A version that uses pointer arithmetic instead of array subscripting :

```
int count_spaces(const char *s) {
    int count = 0;
    for (; *s != '\0'; s++)
        if (*s == ' ') count++;
    return count;
}
```

- Questions raised by the count_spaces example:

- *Is it better to use array operations or pointer operations to access the characters in a string?*
We can use either. Traditionally, C programmers lean toward using pointer operations.
- *Should a string parameter be declared as an array or as a pointer?*
There's no difference between the two.
- *Does the form of the parameter (s[] or *s) affect what can be supplied as an argument?*
No.

Using the C String Library

- Common string operations include copying strings, comparing strings, concatenating strings, selecting substrings.
- Strings are treated as arrays in C, so they're restricted in the same ways as arrays.

In particular, they can't be copied or compared using operators in C.

- Copying a string into a character array using the = operator is not possible:

```
char str1[10], str2[10];
...
str1 = "abc";    /*** WRONG ***/
str2 = str1;    /*** WRONG ***/
```

- Using an array name as the left operand of = is illegal.

Initializing a character array using = is legal:

```
char str1[10] = "abc";
```

- Attempting to compare strings using a relational or equality operator is legal but won't produce the desired result:

```
if(str1 == str2);
```

- This statement compares str1 and str2 as *pointers*.

Since str1 and str2 have different addresses, the expression str1 == str2 must have the value 0.

The strcpy function

- The C library provides a rich set of functions for performing operations on strings.
- Programs that need string operations should contain the following line:

```
#include <string.h>
```

- In subsequent examples, assume that str1 and str2 are character arrays used as strings.

- **String copy:** Prototype for the **strcpy** function:

```
char *strcpy(char *s1, const char *s2);
```

strcpy copies the string s2 into the string s1.

- To be precise, we should say “strcpy copies the string pointed to by s2 into the array pointed to by s1.”

strcpy returns s1 (a pointer to the destination string).

- **EXAMPLE:**

```
...
```

```
strcpy(str2, "abcd");//str2 now has "abcd"
strcpy(str1, str2); // str1 now has "abcd"
```

- In the call strcpy(str1, str2), strcpy has no way to check that the str2 string will fit in the array pointed to by str1. If it doesn't, undefined behavior occurs.

- Calling the **strncpy** function is a safer, albeit slower, way to copy a string.

strncpy has a third argument that limits the number of characters that will be copied.

A call of strncpy that copies str2 into str1:

```
strncpy(str1, str2, sizeof(str1)-1);
str1[sizeof(str1)-1] = '\0';
```

The strlen (String Length) Function

- Prototype for the strlen function:

```
size_t strlen(const char *s);
```

size_t is a typedef name that represents one of C's unsigned integer types.

strlen returns the length of a string s, not including the null character.

- Examples:

```
int len;

len = strlen("abc"); // len is now 3
strcpy(str1, "abc");

len = strlen(str1); // len is now 3
```

The strcat (String Concatenation) Function

- Prototype for the strcat function:

```
char *strcat(char *s1, const char *s2);
```

- strcat appends the contents of the string s2 to the end of the string s1.

- It returns s1 (a pointer to the resulting string).

- examples:

```
strcpy(str1, "abc");
strcat(str1, "def");//has "abcdef"
```

```
strcpy(str1, "abc");
strcpy(str2, "def");
strcat(str1, str2);//has "abcdef"
```

```
strcpy(str1, "abc");
strcpy(str2, "def");
strcat(str1, strcat(str2, "ghi"));
/* str1 has "abcdefghi";
   str2 has "defghi" */
```

- strcat(str1, str2) causes undefined behavior if the str1 array isn't long enough to accommodate the characters from str2.

- Example:

```
char str1[6] = "abc";
strcat(str1, "def"); /** WRONG **/
```

- str1 is limited to six characters, causing strcat to write past the end of the array.

- The **strncat** function is a safer but slower version of strcat.

- Like strncpy, it has a third argument that limits the number of characters it will copy.

- A call of strncat:

```
strncat(str1, str2, sizeof(str1) - strlen(str1) - 1);
```

- strncat will terminate str1 with a null character, which isn't included in the third argument.